

Keeping the machines running

Research-to-submission playbook | The Ken 2026

Prepared 7 September 2026 for a Hyderabad-based PDM student team entering Product Strategy. Official Assembly deadline: 10 September, 11:59 p.m. IST. Aim to submit by 8 p.m. that day. This is a research and execution plan, not a completed customer study or final competition entry.

The strategic recommendation: recruit people who actually coordinate household repairs, reconstruct recent incidents with artifacts, and identify the reason they cannot hand the job over. Design the agent around that reason. Your advantage must come from access and learning, not the number of features or the number of survey responses.

The official brief already proposes household service memory, low-effort collection of appliance information, and persistent service coordination. Repeating those ideas will not satisfy the demand for a new customer insight. Your strongest possible contribution is a discovered mechanism that changes an action, permission, escalation or completion rule.

The three-day operating picture

DAY 1

Evidence

Answers 01 + 02

DAY 2

Design

Loop + rail ask

DAY 3

Submit

Verify + receipt

Day 1 | Evidence: team connection audit, two pilot interviews, targeted recruitment, 6-8 customer interviews, two counterparty interviews, artifact collection, a short branching survey, and first drafts of Answers 1 and 2.

Day 2 | Design: challenge the leading insight with additional users, map existing alternatives, specify the agent's six-step loop and failure paths, and choose one rail capability to investigate deeply.

Day 3 | Submission: test comprehension and permission boundaries, finish all ten answers, curate evidence, check the actual form, and submit with a buffer.

Evidence chain: incident -> artifact -> observed behaviour -> alternative explanations -> customer insight -> product rule -> rail dependency -> test -> submission answer.

What is established and what is still open

Established: the official questions, limits, public deadline, broad judging criteria, selected opening's narrative, and public rail capabilities. Still open: team-specific advantage, actual customer findings, the best segment, final innovation rail, exact logged-in upload constraints and your team registration status. No primary research has been conducted in this report.

Read pages in this order: requirements; first-day plan; recruitment and interviews; survey; synthesis; agent and rails; submission checks. Use the companion fieldwork templates to record genuine work as it happens.

Source: [The Ken: 2026 competition and Assembly](#), retrieved 7 September 2026. Recommendations and research targets throughout are our proposed operating choices unless identified as official requirements.

1. The confirmed submission contract

The public page specifies the following. Labels below are paraphrased; preserve the exact field wording when copying answers into the form. [Official Assembly instructions](#).

01 | Team and connection. One line per member giving name and closest personal connection; a separate team statement of 50 words explaining distinctive access, observation opportunity or prior action. Generic student credentials will not answer this.

02 | One discovered customer insight. 60 words. Evidence is optional but scored; multiple conversations can be attached. The judges read this answer first. The insight must go beyond the opening and visibly change the solution. Accepted examples include timestamped recordings, workaround photographs, a survey with questions and raw responses plus an unexpected result, working notes, and chat logs including AI conversations.

03 | Agent operation. Six cells, one sentence each, maximum 15 words per cell. Cover trigger, prior knowledge, actions, counterparties, human involvement and completion.

04 | Rail roles. One sentence each for payments, logistics and voice. A rail may have no role if you explain why. No numerical word cap is displayed publicly.

05 | Innovation rail. Select one rail and use 40 words to identify the capability you need that does not exist today. A routine integration is insufficient; research and label uncertainty about absence.

06 | Customer asset. 30 words: the one thing the customer hands over and why the exchange is worthwhile.

07 | Adjacent use case. 30 words: what you take over next and why your position gives you an advantage over the current provider.

08 | Personal boundary. Choose one of the sixteen openings you would never delegate, even to a flawless assistant; explain in one sentence. This tests your values, not technical reliability.

09 | Incumbent counterfactual. One Indian company, 60 words: it already has the relevant users, data and payment relationship; explain your best hypothesis for its inaction.

10 | Track. Select Product Strategy. The Assembly-specific instruction says this choice locks on submission. The stated eventual output is a fully specified build, including complete flows, segment map, and the infrastructure limitation encountered and its cause.

Dates and contradictions to handle conservatively

The current official updates list registration closing 8 September; Assembly closing 10 September at 11:59 p.m. IST; shortlist and rails round opening 15 September; finale 10 October; publication late October. The older Unstop listing gives different stage dates and approximately 30 spaces. Use the current official sixteen-opening Assembly page.

Generic eligibility text and the terms allow later track changes, but Q10 says submission locks the track. Treat submission as final unless the organiser confirms otherwise. The picker warns of at most three opening switches; the ground rules also say evidence restarts when switching. Keep the selected title stable.

2. What the opening already tells you

The current picker lists your title second, while its image asset retains the number 11. Select **Keeping the machines running by title**, not its position. This explains the mismatch with your reference to Space 11.

The opening describes household installation, repair and maintenance across appliances and home services. Examples include fragmented water-purifier installation, disagreement between a tenant and landlord, repeated technician follow-ups, and service-contract coordination. Its suggested foundation is an automatically assembled service history, with an agent pursuing work to completion despite closed brand systems. These are the organiser's starting observations, not findings your team can claim to have discovered. [Official opening, embedded in the competition page.](#)

The official narrative describes at least 49 respondents, substantial metro concentration, broad age coverage, and repetition among roughly half of people who tagged the opening. That is a description of this submitted survey material, not an estimate for Indian households. The chart distinguishes self-tagged behaviour from inferred behaviour; it does not provide the denominators needed to rank population prevalence by age.

Audit of your supplied selection playbook

I read the visible text and embedded case/ranking data in `The_Ken_2026_Case_Selection_Playbook.html`. Its ranking combines judgement about recurrence, feasibility, access and differentiation. The weights and scores are a decision aid, not the official judging rubric.

Its stored 18-24 entries for machines are F=4, R=3, P=1, whereas trip planning and cooking have more repetition markers. The current official chart also differs from that stored snapshot. Therefore, neither the file nor the live chart supports claiming that machines have the highest repetitive rate among ages 18-24. Do not combine inferred frustration with self-reported frustration or treat marker counts as unique respondents.

The playbook references `Case studies.docx` and `Pasted markdown(1).md`; neither was present in the workspace. Its underlying extraction cannot be fully reproduced. Its five-case research tournament would be too broad now that you have chosen an opening and have three days.

The useful problem frame

How does a household get a machine back into dependable service without repeatedly supplying context, coordinating other people's commitments, or taking decisions it lacks authority to make?

Start with adults who coordinated an incident, especially accessible young renters. Separate user, owner, payer, approver and person admitting the technician. Age is a recruitment lens; responsibility is the behavioural segment. Do not equate delay with laziness: competing duties, uncertain costs, warranty constraints and disputed responsibility are all explanations to test.

3. What prior competition material teaches

The Ken's official 2025 announcement names IIM Ahmedabad's A-Team as winner. Their concept moved suitable care from hospital to home through AI coordination and episode pricing. The organiser explicitly credited system-level thinking about alternatives, incentives, risks and dependencies. This supports a strategy centred on the whole service episode rather than a catalogue of AI features. [The Ken's winner announcement](#).

The 2025 challenge concerned disrupting an Indian company; 2026 begins with consumer openings and infrastructure rails. Copying last year's company-disruption format would miss the current task. [Zerodha's 2025 announcement](#).

Transfer to your strategy

Map the full episode. Identify who notices the issue, finds a provider, explains the fault, approves the cost, grants access, sources a part, verifies service and handles recurrence.

Explain incentives. A technician may prefer a familiar customer and predictable payment; a landlord may want a repair record; a tenant may want comfort but lack spending authority. These are hypotheses to interview, not statements about every Indian household.

Demonstrate a failure. A compelling agent flow should survive a vague appointment, missing part or revised quote. Show what becomes a recorded state, what action follows, and when a human intervenes.

Check alternatives early. Urban Company's Hyderabad microwave page describes inspection and quote approval, parts sourcing, workshop servicing where needed and post-repair warranty. These capabilities make a generic booking-and-chasing pitch weak. They do not establish that every customer's coordination problem is solved. [Urban Company: Hyderabad microwave repair](#).

What would make this competitive

One memorable insight supported by real incidents; a narrow and defensible initial segment; a loop that removes a demonstrated job; one precise rail limitation; a credible counterparty incentive; and measures that distinguish an appointment booked from a repair completed.

Judge the submission against the official five dimensions: evidence, creativity, clarity, feasibility and thoroughness. There are no published weights on the main page. The round also stresses evidence quality, depth of the rail the solution depends on, and specificity of the requested capability. [2026 judging and Q10](#).

The public showcase links prototype websites. Review of downloaded site content, rather than interactive testing or slide decks, found that ArogyaGhar.ai presents hospital and patient views, a staged care journey, consent, episode economics and exception handling. FinVaani presents a focused vernacular-investor journey using voice and visual guidance. The useful lesson is to make a specific user's job and decision boundaries tangible. Prototype testimonials, partner badges and projected outcomes were not independently verified. [Official 2025 showcase](#); [ArogyaGhar.ai prototype](#); [FinVaani prototype](#).

No public material establishes a formula for winning. This plan improves evidence quality and strategic specificity; it cannot promise selection.

4. First 24 hours: finish defensible drafts of 01 and 02

Assume three teammates. Assign A to research coordination, B to fieldwork and counterparties, C to evidence and product synthesis. Rotate interviewing; each interviewer writes notes before group discussion. If only two people are available, reduce the interview target before compromising consent or evidence quality.

Hours 0-1 | Set up. Confirm all members have completed registration. Copy the ten fields into a draft. Each member writes a personal incident and access inventory independently. Create participant IDs and the evidence folder. Decide what consent covers. Output: Q1 facts, recruitment list, hypothesis log.

Hours 1-3 | Pilot and recruit. Run two 25-minute customer interviews, each followed by a 10-minute debrief. Pilot the survey with three eligible adults using think-aloud questions. Schedule later interviews through at least two recruitment routes. Output: revised guide and survey v1; pilot responses labelled separately.

Hours 3-8 | Reconstruct incidents. Conduct four more customer interviews and two technician/landlord/service-coordinator interviews in parallel where team availability permits. Spend 10 minutes logging evidence immediately after each. Observe one existing workaround with permission. Launch the revised survey. Output: timelines, permission map, artifact records.

Hours 8-10 | Find contradictions. Compare independent notes before agreeing themes. Separate what is said, shown, inferred and unknown. Identify two candidate insights and a rival explanation for each. Recruit a satisfied user or someone who solved the same task without repeated effort.

Hours 10-14 | Close evidence gaps. Conduct up to two targeted customer interviews; obtain missing artifact permission or clarifying timestamps. If a real unresolved case exists and the user agrees, observe the next natural step. Do not create a repair problem or a fake interaction.

Hours 14-17 | Draft and trace. Complete Q1 from verified member facts. Draft Q2 around one mechanism and name the design change it causes. Link every clause to participant and evidence IDs. Draft the corresponding Q3 step now; do not defer this linkage.

Hours 17-20 | Challenge. A teammate argues that the insight is already in the brief or contradicted by the evidence. Revisit the original recording, not only the summary. Ask a relevant participant whether the factual reconstruction is accurate without requesting endorsement of your conclusion.

By hour 24 | Freeze v1. Archive raw evidence, preserve contradictory cases and record the confidence level. Produce Q1 v1, Q2 v1, an evidence index, and a list of uncertainties for Day 2. If evidence remains weak, label Q2 provisional rather than manufacturing a discovery.

Suggested first-day target: 6-8 customers, two counterparties, 4-6 relevant artifacts and 20-30 usable survey responses if recruitment permits. These are scheduling targets, not proof thresholds or organiser requirements. A smaller well-documented study can be stronger than a larger shallow one.

5. Research objectives and recruitment

Five questions the research must answer

R1 | Where does human work accumulate? Count distinct service episodes, contacts within an episode, elapsed downtime and active coordination time separately. This determines whether there is a recurring job worth delegating.

R2 | Why does the workaround persist? Reconstruct what the person is protecting and the tradeoff made at the moment of a decision. This supplies the possible Q2 insight.

R3 | Who has authority? Distinguish the coordinator from the machine owner, payer, access provider and repair approver. This determines permission gates and payment rules.

R4 | What context is missing? Watch the user retrieve records; determine whether data entry, retrieval, trust or access is the obstacle. This determines Q6 and onboarding.

R5 | What prevents completion? Examine successful and failed service promises with counterparties. Determine whether the constraint is information, capacity, incentives, stock, approval, address or transport. This determines Q3-Q5 and test design.

Sampling plan for a three-day sprint

Aim for 10-12 customer interviews total, with overlap recorded rather than disguised as separate samples. Start with 4-5 adults aged 18-24 who personally coordinated a recent repair; 3-4 other household coordinators for comparison; and 2-3 customers whose latest experience was smooth. Include owned and rented appliances, brand and independent service, warranty/AMC and paid repairs where access permits.

Add 3-4 counterparty interviews: an independent technician, a service-centre or parts coordinator, and a landlord/PG or apartment manager. PG operators offer process access but do not prove a household consumer thesis. Keep their findings separate.

Recruit through classmates who actually handled repairs, family networks, nearby rental communities and local repair shops. Record which route produced each participant. Seek variation in living arrangement, language and decision authority; do not collect sensitive demographics without a research purpose.

Neutral recruitment script

"We are students studying how people arrange household machine installation, servicing and repairs. We would like to understand one recent experience, whether easy or difficult. It takes about 25 minutes. Participation is voluntary. We will explain recording, competition submission and possible publication separately before starting. We are not selling a service."

Four screening questions

Are you 18 or older? Have you personally contacted or coordinated someone about a household appliance/service incident in the past 90 days? What was the most recent machine or service? What did you personally do?

Prioritise those with direct actions and recent recall. If no 90-day incident exists, accept a 91-180-day case only as a separately labelled supplementary case. People who merely witnessed a repair are useful for context, not equivalent coordinator evidence. Limit repeat household incidents to avoid treating several relatives as independent confirmation.

6. Consent and evidence collection

The competition requires informed consent for interviews, recording, submission and organiser use, including publication. Submissions must be in English; retain original-language evidence and provide checked English translations. AI help is permitted, but fabricated evidence is explicitly disqualifying. [Competition terms, sections 2 and 4.](#)

Participant briefing: read before research

"We are [names], students at [institution], researching household repair coordination for The Ken's 2026 competition. Participation is voluntary; you can skip questions or stop. We would like to take notes and, if you agree, record this conversation. With your permission we may submit your words and selected evidence to The Ken and its judges and partners. Material may be excerpted, adapted or used in published work; the competition terms allow broad ongoing use. The public instructions describe quotes by age band and city, without names, but a voice or photograph may still identify someone. We will remove unnecessary personal details. We cannot promise removal after submission or publication. Please tell us what you are comfortable sharing."

Ask separately and record yes/no: participation and notes; audio; video if actually needed; named artifacts; submission to the organiser; the organiser's stated reuse/publication terms. Provide the terms link and time for questions. Ask the participant to explain in their own words where the material could appear. If submission/reuse permission is declined, keep the session out of the submitted evidence and claims unless appropriate consent is subsequently obtained. Recruit adults for this sprint.

For each evidence type: method and output

Recording: critical-incident interview. Obtain consent before recording, test audio, keep the original, record the exact timestamp of the relevant passage and 30-60 seconds of context. Produce an English excerpt with translation reviewer and original timestamp. Do not splice a new meaning.

Workaround photo: contextual observation. Ask the participant to show how they retrieve a service number, invoice or history. Photograph the actual system with permission and caption its purpose, date and observed use. A photo of a broken machine proves its condition only, not the coordination mechanism. Never stage a notebook as pre-existing.

Survey: use the versioned questionnaire on the following pages. Preserve questions, routing, raw responses, recruitment channels, exclusions and one genuinely unexpected answer. Keep optional follow-up contact details in a separate file.

Working notes: timestamped decision record. Capture assumptions before interviews, observations immediately afterward, alternative explanations, discarded ideas and the reason for each product change. A newly drawn journey is a reconstruction; label it with the interview source and drawing date.

Chat logs: later export relevant threads with permission, enough context and timestamps; redact third-party identifiers. Retain AI prompts and outputs as process evidence, clearly labelled AI-assisted. An AI-generated persona or conversation is never customer evidence.

Use originals privately; create redacted submission copies. Do not upload entire phone libraries, contact books, identity documents, payment credentials or unrelated household chats. The organiser's reuse licence makes careful selection materially relevant.

7. Customer interview: exact guide

Allow 25-35 minutes plus a 10-minute debrief. Ask the main question first, then only the probes needed. Do not mention an agent until the final delegation section. Use the participant's preferred language; have a fluent teammate review translations.

1. **"Tell me about the most recent time you arranged installation, service or repair for a machine at home."** Probe: approximately when; which machine; where you lived; your role. Purpose: one bounded incident and eligibility, not a general complaint.
2. **"What happened first?"** Probe: what alerted you; what you did next; who else acted. Repeat "And then?" until the story ends. Purpose: establish sequence before imposing categories.
3. **"Could you walk me through the contacts you made?"** Probe: first contact, later calls/messages, callbacks, people involved and channel changes. Ask permission to compare with logs. Purpose: distinguish contact volume from repair frequency.
4. **"What information did you need at that point?"** Probe: where it was kept; what could not be found; what had to be repeated; what happened because it was missing. If willing, ask "Show me how you would find it now." Purpose: observe the real data and memory burden.
5. **"How did you decide whom to contact?"** Probe: alternatives considered, previously used provider, warranty or relationship considerations, why a service app was or was not used. Purpose: find switching constraints and existing substitutes.
6. **"What did the other person commit to?"** Probe: exact words recalled, a time window versus vague promise, who confirmed, what actually happened. Ask for the corresponding artifact if available. Purpose: distinguish no promise from a broken commitment.
7. **"Tell me about any point where the next action was unclear."** Probe: who could decide; what information or approval was missing; what happened if you waited. If everything was clear, record that. Purpose: locate ownership and authority gaps without suggesting them.
8. **"How was the cost decided?"** Probe separately: who owned the machine; who approved; who paid; whether the quote changed; deposits, reimbursement or warranty checks. Purpose: reconstruct permissions and payment events.
9. **"Did any item need to move somewhere?"** Probe only if yes: what item, who chose the part, who packed it, who transported it, handover record, delays and return. If no, record no movement. Purpose: test whether logistics is necessary.
10. **"When was the machine usable again?"** Probe: how they knew; what remained unresolved; repeat fault; warranty record. Purpose: define completion independently of a provider's status update.
11. **"What did this require from you?"** Probe separately: minutes actively coordinating; time waiting at home; days without the machine; actual expenses; consequences. Ask what is remembered versus documented. Purpose: measurable baseline without conflating time categories.

8. Deep probes, trust and counterparty interviews

Continue the customer guide only as time permits.

12. "Which moment was hardest, if any?" Then "What mattered to you at that point?" and "Can you describe what you did because of that?" Purpose: uncover protected outcomes through behaviour. Avoid repeatedly asking accusatory "Why didn't you...?"

13. "Was there a moment you postponed the next step?" If yes: "What was happening then?" "What needed to change before you acted?" "What happened the time before?" Purpose: distinguish deferral, competing priorities and dependency. Never label it laziness.

14. "Tell me about another repair that went differently." Probe: what made it easier or harder; same person or different; what differed in authority, parts and provider. Purpose: rival explanation and negative case.

15. "If a person you trusted had handled this incident, which actions could they take without checking with you?" Ask separately about booking, sharing selected history, changing provider, approving a revised quote, paying, arranging collection and confirming completion. Ask what they must bring back for approval. Purpose: delegation boundary before AI preferences.

16. "What would make you take the job back from them?" Then "What record would reassure you it was finished?" End: "What have I misunderstood or not asked?" Purpose: recovery, verification and overlooked explanations.

Technician or service-centre guide: 20-25 minutes

1. "Walk me through the last household job from first contact to closure." Identify work type and personal role; do not request identifiable customer records.

2. "How do you decide which job happens next?" Probe travel, urgency, parts, payment and existing promises only after an open answer. This tests scheduling incentives.

3. "Describe the last appointment that changed." Probe the cause, who informed whom, what notice was possible and how it was resolved. Avoid presuming dishonesty.

4. "What information is most often missing before a visit?" Ask for a de-identified job card or blank form. This tests structured service memory.

5. "When can you quote accurately?" Probe inspection, part compatibility, who sources parts and what causes revisions. This bounds agent diagnosis and purchasing.

6. "Who is allowed to authorise work and payment?" Ask about an absent owner, tenant or unknown assistant. This tests counterparty recognition of delegation.

7. "What proves a job is complete, and what happens if the fault returns?" Probe job cards, customer checks, receipts and rework incentives.

8. "What would make coordination by an assistant easier or harder?" Ask at the end; distinguish hypothetical willingness from demonstrated operational acceptance.

For landlords/PG managers add: "Describe the last disagreement about a repair cost"; "What evidence did you need to approve it?"; "Who admitted the technician?"; "What was the last repair a tenant could approve themselves?" Do not assume the tenant can transfer the owner's authority.

9. Survey design and analysis contract

Use a short core with branching, followed by an optional incident module. A detailed design does not require every respondent to complete a long questionnaire. Target 7-9 minutes for the relevant path; measure this in pilots rather than promise it untested.

Purpose: describe recent coordinator experiences in your accessible sample, recruit interviews and identify variation. The survey cannot independently establish the hidden mechanism, Indian market prevalence, causal impact or product demand.

Design rules

Ask about one recent incident before opinions. Ask one concept at a time. Offer zero, none, not applicable and cannot recall where meaningful. Keep ordinal scales ordered; rotate only unordered answer lists when the survey platform supports it, keeping Other/None last. Show open questions before lists that could prime them.

Pew's questionnaire guidance supports neutral wording, avoiding multiple concepts in a single question, considering question-order effects and pretesting. The instrument here is designed for your task, not copied or represented as a validated scale. [Pew: Writing survey questions](#).

Before launch

Run three cognitive pilots with eligible adults. Ask: "What did this question mean to you?" "How did you recall that answer?" "Was an answer option missing?" Test every skip path. Have a second fluent speaker check translated versions against the intended meaning. Log changes; if wording changes materially, exclude pilot responses from the main analysis or report them separately.

Administration

Record survey version, date, language, recruitment channel and anonymous response ID. Do not require email login solely for deduplication. Offer a separate follow-up form so contact information is not mixed with raw answer exports. Do not solicit only frustrated customers or reward severe answers. If compensation is offered, use the same terms regardless of the story and disclose it.

Analysis decided before seeing results

Primary unit: one respondent's latest eligible incident. Household duplicates and interview overlap must be flagged. Separate unresolved cases from completed cases. Distinguish missing answers from zero and not applicable. Report counts with denominators, for example "5 of 9 respondents who answered Q14"; do not invent an actual result from this example.

Summarise incident counts and follow-ups separately, plus median active effort among respondents who can recall it. For open answers, code independently before group synthesis. Preserve minority and contradictory cases. Cross-tab authority role, channel and living arrangement only when cell sizes justify a readable description; avoid spurious percentages from tiny cells.

You may aim for 30-50 usable responses across the sprint. Recruitment is convenience-based and likely Hyderabad/student-network heavy; report that limitation. Do not use a margin of error or call the sample representative. Attach the exact questionnaire, cleaned analysis and redacted raw responses, with an exclusion log.

10. Survey: screening and incident facts

Intro: explain voluntary participation, study purpose and the consent terms in section 6. Display research contact [team-controlled contact] and intended submission date. No participant data should be collected until the team replaces these placeholders.

S0. "Do you agree to participate under the explanation above, including submission and the stated organiser use of your anonymised answers?" Yes / No. No ends the survey. If the participant does not understand, resolve questions before proceeding. No recording or photographs are requested by this survey consent.

S1. "Are you 18 or older?" Yes / No / Prefer not to say. Only Yes proceeds for this sprint.

S2. "In the past 90 days, have you personally helped arrange installation, servicing or repair of a household machine?" Yes / No / Cannot recall. Define arrange as contacting a provider, coordinating a visit, obtaining approval, arranging parts or handling payment. If No/Cannot recall, skip to S22-S23 and the end; do not ask them to imagine an incident.

S3. "How many separate installation, service or repair episodes did you help arrange in those 90 days?" 1 / 2 / 3 / 4 / 5 or more / Cannot recall. Explain: several calls about the same unresolved issue count as one episode. Link: R1, recurrence.

"For the next questions, think only about the most recent episode."

S4. "What machine was involved?" AC / Refrigerator / Washing machine / Water purifier / Geyser / Kitchen appliance / Other [text]. One answer; if several, choose the most recently addressed. Link: segment.

S5. "What was the main reason for arranging service?" Installation / Scheduled maintenance / Fault or breakdown / Follow-up on a previous repair / Other [text] / Cannot recall. Link: trigger.

S6. "Approximately when did this episode begin?" Within 7 days / 8-30 days ago / 31-60 / 61-90 / Cannot recall. Link: recall reliability.

S7. "Who owned this machine?" Me / Someone else in my household / Landlord or PG operator / Shared community or building / Other [text] / Do not know. Link: authority.

S8. "Which actions did you personally take?" Select all: contacted provider; arranged visit; supplied records; approved work; paid provider; arranged parts or transport; admitted provider; checked completion; other [text]. Link: respondent role. If no direct action, flag for eligibility review.

S9. "Which channel did you use first to arrange this episode?" Brand service centre / Independent technician / Service marketplace such as Urban Company / Landlord or PG manager / Building staff / Retailer / Other [text] / Cannot recall. Link: substitute.

S10. "What happened after that first contact?" Optional short text. Link: unprimed sequence and unexpected discovery; ask before the structured effort questions.

11. Survey: effort, outcome and segmentation

S11. "After the first contact, how many follow-up contacts did you initiate about this same episode?" 0 / 1 / 2 / 3-4 / 5-7 / 8 or more / Cannot recall. Count distinct call attempts or message exchanges, not every message bubble. Link: coordination burden.

S12. "About how many minutes did you actively spend arranging this episode in total?" Numeric minutes / Cannot recall. Include calling, messaging, finding records and approvals; exclude passive waiting and hands-on repair. Link: effort baseline.

S13. "What is the episode's current status?" Completed and machine usable / Provider finished but issue remains / Still waiting for service or repair / Cancelled or abandoned / No service needed after assessment / Other [text]. Link: outcome, avoids equating booking with completion.

S14. If S13 is completed: "About how long passed between the first request and the machine becoming usable?" Same day / 1-2 days / 3-7 days / 8-14 days / 15 or more days / Cannot recall. Link: restoration time. Unresolved episodes do not receive a fabricated completion time.

S15. "Overall, how frustrating was arranging this episode?" Not at all / Slightly / Moderately / Very / Extremely / Prefer not to answer. Link: reported experience; not a diagnostic scale.

S16. "Was there a time you decided to leave the next action until later?" Yes / No / Cannot recall. If Yes: **S16a. "What was the main reason at that time?"** Optional free text. Link: R2; no suggested laziness/procrastination label.

S17. "Who had final authority to approve repair spending?" Me / Another household member / Landlord or PG manager / Building or association representative / Joint approval / Other [text] / No spending decision was needed / Do not know. Link: permission design.

S18. "Who paid the provider?" Me / Another household member / Landlord or PG operator / Building or association / Other [text] / Not paid yet / No payment required / Do not know. Link: payer, distinct from approval.

S19. "Did a physical item need to be transported for this episode?" Yes / No / Do not know. If Yes:

S19a. "What moved?" Part or consumable / Whole machine / Other [text]; select all. **S19b. "Who arranged that movement?"** Me / Provider / Another person [role] / Do not know. Link: logistics necessity.

S20. "What, if anything, did you do to make the arrangements easier?" Optional text. **S21. "What was the most important thing you wanted to avoid during this episode?"** Optional text. Link: workaround and protected outcome; investigate through interviews before interpreting causally.

S22. "What is your age band?" 18-24 / 25-34 / 35-44 / 45-54 / 55-64 / 65+ / Prefer not to say. **S23. "Which city do you currently live in?"** Text / Prefer not to say. Optional **S23a: living arrangement:** family home / rented alone / shared rental / PG-hostel / other / prefer not. Link: scope disclosure.

12. Optional survey module and observation tests

Invite only S2=Yes respondents: "Would you answer a few more questions about the same episode?" Yes / No. No goes to the separate optional interview invitation.

S24. "Was warranty or an AMC relevant to this episode?" Yes / No / Unsure. If Yes: **S24a. "How did you find the relevant information?"** Paper document / Email / Message or photo / Provider told me / Someone else found it / Could not find it / Other [text]. Link: customer asset and retrieval.

S25. "Did the quoted amount change after you first agreed to it?" Yes / No / No quote agreed / Cannot recall. If Yes: **S25a. "What happened next?"** Text. Link: changed scope and reauthorisation.

S26. "About how much did your household actually pay for this episode so far?" Optional numeric INR / Nothing / Do not know / Prefer not to say. **S27. "Did you pay anyone separately to handle the arrangements?"** Yes / No / Do not know. If Yes, optional amount. Link: actual spend and existing outsourcing, not willingness-to-pay fiction.

S28. "Which one action would you most want another person to handle next time?" Free text / None. **S29. "Which one decision would you keep for yourself?"** Free text / None / Unsure. Link: delegation contrast. These are stated preferences, not demonstrated adoption.

S30. "Is there anything about this experience we have misunderstood or missed?" Optional text. End with a separate follow-up link asking whether they are willing to discuss the incident; do not auto-enrol anyone.

Contextual observation: 15 minutes

Ask a consenting participant to retrieve the information they actually used for the most recent service. Let them choose the device and record only the relevant steps. Note time, switches between locations/apps, assistance needed, missing records and any mismatch between what they said and did. Do not ask them to expose unrelated messages. A failed retrieval is evidence; do not repair the task before observing it.

One-day event diary: only for existing live cases

At each naturally occurring update ask the participant to record time, who initiated, event, next action, responsible person, promised time, effort and optional artifact. At day's end ask what remains unresolved. Label the short observation window; it cannot establish long-term recurrence.

Day 2 concierge test

With explicit participant and counterparty permission, a teammate manually performs one bounded coordination task as a student researcher. Record the user's allowed actions and approval gates first. Stop at spending, home access or changed scope unless separately authorised. Compare active user effort with the reconstructed baseline, but do not claim causal time savings from unlike incidents. A human-assisted test demonstrates workflow plausibility, not sandbox integration or autonomous reliability.

13. Working notes that show real thinking

Before interviews: 35-minute ideation session

First 5 minutes: each member writes what they have personally experienced. Next 10: map a generic service journey using a distinct colour for assumptions. Next 10: list competing explanations and what would disprove each. Final 10: select the riskiest unknowns and assign interviews. Photograph the actual work with its date; preserve discarded ideas. Do not start by voting for AI features.

Initial hypotheses: test all, commit to none

H1: users keep a familiar provider despite delays because replacement threatens warranty, quality or relationship. Disconfirm with people who switch easily. Possible design change: preserve provider preference and ask before switching.

H2: an apparent service delay is really an approval delay between coordinator, owner and payer. Disconfirm with sole-owner incidents that show the same mechanism. Possible change: establish a role-specific approval packet before scheduling.

H3: repeat visits happen because model or part information is missing before the first visit. Disconfirm where complete information was available but capacity or stock caused delay. Possible change: technician-confirmed job readiness before dispatch.

H4: customers cannot tell when to intervene because promises have no reliable next checkpoint. Disconfirm where confirmed appointments still fail for reasons no information can fix. Possible change: record uncertainty and escalation thresholds rather than automatically treating every promise as firm.

H5: creating appliance records is abandoned because people do not receive immediate value. Disconfirm with users who maintain records successfully. Possible change: build memory incident by incident, starting with one repair packet.

After every interview: ten-minute independent note

Record participant ID, role, date, interviewer, episode, facts, artifact IDs, exact quote and timestamp, interpretation, rival explanation, surprise, design implication and follow-up. Mark every item **said / observed / artifact-supported / inferred / unknown**. A screenshot and a quote from the same incident corroborate it; they are not two independent people.

Synthesis session: 60-90 minutes

Lay out incidents, not feature suggestions. Group by mechanism: fragmented context, authority, uncertain scheduling, parts readiness, warranty, access, payment, closure or no significant pain. Keep emergent codes that do not fit this list. Compare roles and outcomes. Select the insight that is new relative to the brief, changes a product rule and survives the strongest contrary case.

Use an evidence chain such as P03 -> E07 at 08:42 -> O04 -> I02 -> D03 -> T02 -> Q2/Q3/Q5. These IDs are examples of notation, not collected evidence. Prefer three independent relevant incidents plus a concrete artifact for a central mechanism when feasible, but no fixed count proves validity. With only one case, state the narrow scope honestly.

14. Answers 01 and 02: drafting and evidence gates

Answer 01: connection becomes a research advantage

Ask every member: Which incident did you personally handle? When? What did you do? What artifact can you show? Whom can you reach today? Why will that person talk candidly? What relevant skill have you demonstrated? What did you learn that changed your view? What access have you actually secured rather than merely hoped for?

Member-line scaffold: **[Name] - [specific personal role in a real incident or recurring responsibility]**. Do not inflate attendance, family connections or student status into expertise.

Team scaffold, under 50 words: **We can observe [specific setting] through [verified access]. [Completed action] gives us [specific learning advantage]. Our team combines [demonstrated capability] with [demonstrated capability] to investigate [narrow uncertainty]**. Fill only with verified facts; this is not submission-ready text.

Evidence can include a dated incident artifact, a permission-backed observation visit, or a genuine prior project. The public instructions explicitly associate attachment scoring with Q2; they do not state a separate compulsory Q1 evidence upload. Maintain Q1 provenance anyway and attach only where the form permits.

Answer 02: one human mechanism and its design consequence

Scaffold, under 60 words: **Among [specific people], we found [non-obvious behaviour]. They were protecting [outcome], so [mechanism behind the behaviour]. [Brief evidence scope] supports this. Therefore our agent [specific action or approval rule], rather than assuming [disproved assumption]**. Rewrite naturally in the team's voice and count words.

Do not use the scaffold to fill in a fashionable insight before interviews. "Customers hate chasing technicians" and "households need service memory" already appear in the brief. A genuine discovery might explain when a customer refuses an apparently faster option, but that must emerge from actual behaviour.

Q2 evidence packet

1. A short evidence index identifying the main claim, participant roles, relevant timestamps and artifact filenames.
2. The strongest consenting interview passage with enough surrounding context, plus its English transcript.
3. A real workaround or incident artifact and a caption explaining precisely what it shows.
4. Cross-case working notes, including the strongest contrary case and why the final claim is bounded.
5. If used, survey questions, routing, redacted raw answers, exclusions, counts and the unexpected answer.

More attachments are permitted, but make the decisive moment easy to find. Do not bury the claim in hours of unindexed footage. A sharp quote does not establish prevalence; a survey percentage does not explain motive; a team note documents interpretation rather than independently proving it.

Day 1 acceptance gate: real member facts; one scoped customer mechanism; at least one traceable supporting episode; documented submission consent; an explicit changed product step; no invented count, quote or outcome. If any gate fails, keep that clause provisional.

15. Derive the agent from evidence

Use a decision record for each proposed feature: supporting incident IDs; observed mechanism; rival explanation; intended user outcome; smallest product behaviour; needed data/authority; rail; failure path; acceptance test. Delete features without evidence or a clearly labelled feasibility necessity.

Conditional product directions

If **approval fragmentation** dominates, specify a service episode with separate coordinator, owner and payer permissions. Test whether this reduces repeated approval requests without letting a tenant authorise the owner's money.

If **trusted-provider retention** dominates, give users provider-preserving recovery and a clear switching threshold. Test with a scenario where a faster unfamiliar provider is available.

If **parts readiness** dominates, require a qualified technician to confirm model and part compatibility before purchase or shipment. Test wrong-part and unavailable-stock paths. The agent is not the repair expert.

If **service-memory capture** dominates, ask for one incident packet and show immediate retrieval value before requesting a wider archive. Test extraction accuracy and whether the participant corrects the wrong model, date or warranty.

Provisional six-step loop: architecture example, not a customer finding

1. A fault report or service date opens an episode for the household's identified machine.
2. The agent retrieves service history, ownership, warranty evidence, contacts and existing permissions.
3. It calls the provider to establish requirements, quote, appointment and the next confirmed checkpoint.
4. It requests missing approvals and coordinates authorised payment or necessary parts movement.
5. It monitors commitments and escalates exceptions within the household's agreed rules.
6. It records completion after customer confirmation, then updates history and unresolved issues.

Each sentence is at most 15 words; modify the steps after selecting the insight. A complete specification must separately cover the branches compressed out of these six cells.

Technical boundary

Use a language model to interpret selected records and conversation, with confidence checks. Keep state transitions, authority, spend checks and retries in explicit rules and an auditable event log. Counterparty speech is evidence to interpret, never permission to expand spending authority. Use transaction identifiers and status checks to avoid duplicate actions.

Do not infer a technical diagnosis, warranty eligibility or part compatibility from uncertain text and then spend money. Escalate unknowns to a qualified provider or authorised human. Preserve the distinction between appointment accepted, technician visited, payment completed, machine usable and customer-confirmed closure.

16. Rails: what is supported and what must be tested

The following separates public capability from your proposed use. Documentation access is not proof of working integration or guaranteed competition sandbox access.

Voice | Gnani

Gnani's public platform describes multilingual voice agents, custom API support and inbound/outbound use cases. This supports investigating provider calls in Telugu, Hindi or English as appropriate. It does not verify the exact competition API, allowed dialling, IVR navigation, recording access or all language performance. [Gnani voice platform](#); [Gnani outbound use cases](#).

Test: consented test calls; mixed-language dates; vague promises; interruptions; background noise; wrong number; refusal to speak to an assistant. Require confirmation of critical extracted details. Unknown or low-confidence commitments stay unconfirmed; cap retries and offer human handoff.

Payments and authorisation | Pine Labs

P3P documents delegated agent payments using mandates and scoped payment tokens, with receipts. Its SDK documentation describes Grantex consent, scopes, transaction limits and budget handling. Therefore, "let an agent pay under a cap" is already a documented capability, not a defensible missing-capability claim. [P3P overview](#); [P3P SDKs](#).

Test: changed quote, revoked or expired authority, pending payment, duplicate retry, actual payer identity and non-integrated merchant. A payment receipt proves payment, not repair quality. Do not assume escrow, automatic refunds, milestone release or support for every local technician merely because payment APIs exist.

Logistics | Delhivery

The Maps developer portal exposes address, geocoding, route and distance tools through MCP. These help location planning; they do not by themselves book a courier or guarantee a technician arrives. [Delhivery Maps developer portal](#).

Separately, Delhivery's client portal describes shipment creation, tracking, pickup requests, serviceability and shipping-cost APIs; access and shipment eligibility must be checked. [Delhivery client developer portal](#).

Test physical movement only where interviews show a part or eligible machine actually travels. Verify packaging, size/weight, serviceability, pickup address, custody and handling restrictions. Never assume a courier diagnoses repairs, installs parts or handles bulky appliances under ordinary parcel terms.

Q4 drafting rule: describe the role actually necessary for your selected episode. All three may be useful across the product, while a particular episode needs only two. If no defensible logistics job exists, explain that candidly as the official instructions allow.

17. Choose one innovation rail and one customer asset

Rail choice: decision after discovery

Score each candidate from 0-2 on four questions: is it central to the observed mechanism; is the missing capability specific; does existing documentation fail to cover it; can you explain a credible request and test? Any candidate with no evidence connection should be dropped regardless of score. This is our decision aid, not the judging rubric.

Candidate A | Voice. A mutually confirmed, machine-readable service commitment with confidence, accountable counterpart, revision history and an escalation handoff. This may partly be application logic; verify what the voice rail truly needs to expose. Strong only if ambiguous commitments cause the observed failure.

Candidate B | Payments/authorisation. A repair-specific authority handoff across coordinator, owner and payer, tied to a revised scope and verifiable service checkpoint. Separate this from already available caps and grants. Strong only if current primitives cannot express a necessary constraint. Do not claim this capability is globally absent without checking the protocol and partner response.

Candidate C | Logistics. A repair handover that binds model/part identity, custody and failed-pickup recovery across household and service centre. Determine which elements are possible with existing shipment metadata and which require a new rail guarantee. Strong only if transport is the observed bottleneck.

Provisional preference: investigate Voice first if interviews show unresolved promise interpretation; investigate Payments/authorisation first if the person coordinating lacks the payer's authority. Do not select a rail solely because its prize is attractive.

A precise rail ask has five parts

Current supported primitive; observed customer failure; missing contract or state; requested capability; acceptance test and fallback. In Q5 compress these into the required 40 words. For strategy-track depth, retain the expanded version and the documentary basis for any claimed limitation.

Customer asset for Q6

A plausible first asset is the **selected service-history packet for one machine**: invoice, warranty/AMC and relevant service thread, supplied by the customer. The exchange is immediate reduction in repeated explanation and retrieval. This is a hypothesis to test, not a reason to request the whole inbox or WhatsApp account.

Ask the participant to choose a real redacted item, watch whether they can retrieve it and ask what they refuse to share. Measure time to usable packet, extraction errors and corrections. If they cannot or will not supply it, change onboarding. Permissions to make calls, arrange visits and spend are separate from permission to read a document.

For Q7, a possible adjacency is renewal or warranty follow-through using the service history already created. Validate retention and demand first. Avoid jumping from one repair to an all-purpose household super-agent without a shared asset, customer and execution advantage.

18. Product strategy, tests and the remaining answers

Business model to investigate, not assert

Compare household payment per episode, an annual multi-appliance plan and distribution through rental/PG or housing managers. Interview the actual buyer for each. PG evidence cannot silently become a consumer subscription forecast. A technician-funded model may create a conflict if the agent claims to represent the household; disclose who pays and what influences provider selection.

Estimate contribution per resolved episode as customer revenue minus voice/API, human handling, support, acquisition allocation and any logistics subsidy or dispute cost. Obtain actual quotes or public prices before assigning values. Test low-frequency households separately from frequent coordinators; repeated calls within one breakdown do not prove monthly retention.

Measure user active coordination minutes; completed episodes confirmed by customer; time to usable machine; human escalations; incorrect actions; repeated faults; and cost per completed episode. Track unapproved spend and incorrect completion as guardrails. Do not count reminders sent or calls placed as customer value.

Three small Day 2 tests

Comprehension: show the six-step flow to 3-5 relevant users. Ask them to explain what happens, who approves, and how it finishes. Record misunderstandings rather than rating attractiveness.

Permission: replay their actual incident with a revised quote, unfamiliar provider or changed appointment. Ask what the agent may do and why. Compare with the proposed policy; change it where it oversteps or asks too often.

Exception tabletop: a teammate plays a technician using an explicitly simulated script. Test no answer, no-show, missing part, price change, pending payment and failed pickup. Label every simulation. For Product Strategy, this validates specified behaviour, not actual rail execution.

Answer 08: elicit authentic values

Each member independently chooses an opening and explains the human value they would retain even with flawless execution. Discuss disagreements and select the team's honest view. Do not answer "health, because AI makes errors"; the premise removes errors. Do not let AI manufacture this belief for the team.

Answer 09: investigate a company with the actual relationship

Urban Company is a candidate because it already offers appliance services and connected repair steps. Its existing capabilities are a reason to investigate the residual gap, not proof it should build your eventual concept. Verify it has the exact users, data and payment relationship your selected segment needs. Then label your explanation of inaction as a hypothesis: for example, incentives might favour jobs inside its network over household coordination across brands and existing technicians. Do not state internal company motives as fact.

Q10 remains Product Strategy. Prioritise a specified flow, boundaries, segment logic, evidence and a precise rail request over a rushed production build.

19. Days 2 and 3: submission operations

Day 2 | Validate and specify

Morning: conduct 3-4 targeted follow-up interviews, including the strongest counterexample and a counterparty relevant to the proposed rail. Check existing alternatives against the same incident. Freeze the scope and insight only after this comparison.

Afternoon: create the actor/permission map, episode state flow, six-step draft, customer asset exchange and rail capability matrix. Run comprehension and exception tests. Record what changed and why. Draft Q4-Q7 and collect the team's independent Q8 views.

Evening: write the company counterfactual, simple unit-economic model with unknowns clearly marked, and evidence-backed outcome measures. Complete a first full ten-answer draft. Have one teammate read only Q2 and Q3 to check that the promised insight actually changes the loop.

Day 3 | Reduce, verify and submit

Morning: close only material evidence gaps. Obtain missing consent, check translations and source timestamps, remove unsupported claims and reconcile numerical denominators. Freeze the main narrative by noon.

Afternoon: inspect the logged-in form for file types, maximum sizes, link permissions, number of attachments and mandatory fields. These constraints were not exposed in the public guidance and remain unverified. Prepare files accordingly. Keep the original evidence secure and submit selected, consented copies.

By 6 p.m. IST: conduct a cold read by someone unfamiliar with the case. Ask them to identify customer, insight, changed step, human boundary, necessary rail and missing capability. If they cannot, simplify the answer instead of adding slides.

By 8 p.m. IST on 10 September: have the captain submit, then verify receipt or success status and save a copy of the submitted fields and evidence manifest. Recheck the official deadline that morning. Do not rely on autosave as proof of submission.

Final checks tied to actual rules

All members registered and same-institution eligible; selected title correct; one team and one opening; team consent to Product Strategy lock; all answers in English; Q1 team statement ≤ 50 words; Q2 ≤ 60 ; six Q3 cells each ≤ 15 ; Q4 three sentences; Q5 ≤ 40 ; Q6 and Q7 ≤ 30 each; Q8 one opening and one sentence; Q9 ≤ 60 .

Evidence links open for the intended reviewers; media plays; decisive passages have timestamps; each asset has documented submission/reuse consent; raw survey answers match reported counts; AI assistance is disclosed wherever the form requests it. Remove placeholders and all hypothetical customer quotes. Preserve a final local copy and receipt.

The terms give the organiser broad reuse rights while leaving ownership with the team and also contain restrictions on external commercial use/showcasing. Review those clauses as a team before later reuse of the submission. This plan does not publish or submit anything on your behalf. [Competition terms](#).

20. Evidence handoff and unresolved decisions

Send this first so we can write Answer 01

For each teammate: name; institution and programme; one real personal connection to this pain; incident date; personal actions; existing artifact; people or places they can access within 24 hours; relevant demonstrated research, design or technical capability; hours available today.

For the team: registration/invitation completion status; who can interview in Telugu/Hindi/English; who can reach an actual repair coordinator or technician; access already agreed; any earlier work on this specific problem; fieldwork budget and practical travel limits.

Send each interview in this format

Participant ID and age band/city; their coordinator/owner/payer role; incident date and machine; recruitment channel; consent status; recording or notes; strongest timestamped passage; relevant redacted artifact; your unexpected observation; alternative explanation; unresolved follow-up. Use pseudonymous IDs in shared research notes and store direct contact details separately.

I can then help code the evidence, challenge the mechanism, identify contradictions, draft a truthful 60-word insight and propagate its implications into Q3-Q9. I cannot replace the team's real customer conversations or author personal beliefs on its behalf.

Current gaps and next actions

Team advantage: unknown; complete the member inventory. **Primary insight:** unknown; run incident interviews. **All three rail roles:** conditional; record actual approvals, calls and item movement.

Innovation claim: unverified; compare a precise ask against current docs and later office hours.

Sandbox access: no session tested; the official page places rails access after shortlisting. **Upload constraints:** captain to inspect logged-in form. **Original playbook sources:** two referenced files missing; supply if available, but do not delay first interviews.

Research review and source notes

The main Ken pages initially failed through the web reader; the public HTML was subsequently retrieved directly. I reviewed the 2026 Assembly text, the embedded selected opening, competition-specific terms, the supplied HTML's visible text and embedded rankings, accessible official rail material, current Urban Company service information and primary 2025 announcements. No registered-team session was entered or changed.

Searches prioritised exact competition requirements, the selected opening, current P3P/Grantex, Delhivery Maps and client APIs, Gnani voice capabilities, a direct alternative and survey design. Broad searches stopped after the official brief was recovered. Remaining gaps require team facts, primary fieldwork, exact partner access or a logged-in form, rather than more generic web research.

The supporting sources are linked beside their claims. Source publication dates are not always available; all public sources were reviewed on 7 September 2026. Website capabilities and dates should be rechecked before final submission. The report's sample targets, hypotheses, interview instruments, gates and strategy options are proposed methods, not observed results.